

Model Optimization and Tuning Phase Template

Date	15 july 2024
Team ID	SWTID1720093035
Project Title	TechPart Vision: Personal Computer Parts Image Classification Using EfficientNet Transfer Learning
Maximum Marks	10 Marks

Model Optimization and Tuning Phase

The Model Optimization and Tuning Phase involves refining neural network models for peak performance. It includes optimized model code, fine-tuning hyperparameters, comparing performance metrics, and justifying the final model selection for enhanced predictive accuracy and efficiency.

Hyperparameter Tuning Documentation (8 Marks):

Model	Tuned Hyperparameters
CNN	<p>Convolutional Layers</p> <ul style="list-style-type: none"> • Number of Filters: 32 • Kernel Size: 3 (first layer), 2 (second layer) • Activation Function: ReLU • Padding: Same <p>Pooling Layers</p> <ul style="list-style-type: none"> • Pooling Size: 2 <p>Regularization</p>

- **Dropout Rate:** 0.5

Dense Layers

- **Units:** 128
- **Kernel Initializer:** HeNormal

Output Layer

- **Units:** 14

Training Parameters

- **Batch Size:** 32
- **Number of Epochs:** 10

```
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout, Flatten
from keras.layers import Convolution2D, MaxPooling2D
import tensorflow as tf

initializer = tf.keras.initializers.HeNormal()

model=Sequential()
model.add(Convolution2D(filters=32, kernel_size=3, padding='same', activation="relu",
                        input_shape=(255, 255, 3)))
model.add(MaxPooling2D(strides=2, pool_size=2, padding="valid"))
model.add(Convolution2D(filters=32, kernel_size=2, padding='same', activation="relu"))
model.add(MaxPooling2D(strides=2, pool_size=2, padding="valid"))
model.add(Dropout(0.5))
model.add(Flatten())
model.add(Dense(128, activation='relu', kernel_initializer=initializer))
model.add(Dropout(0.5))
model.add(Dense(14, activation='softmax'))
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(train_gen, validation_data= val_gen, batch_size= 32, epochs = 10, verbose = 1)
```

<p>EfficientNetV2 B1 Model Initialisation</p>	<p>Tuned Hyperparameters:</p> <p>Base Model: EfficientNetV2B1</p> <p>Base Model</p> <ul style="list-style-type: none"> • Base Model: EfficientNetV2B1 (pretrained on ImageNet, used as a feature extractor) <p>Pooling</p> <ul style="list-style-type: none"> • Pooling: Max (global max pooling applied to the output of the base model) <p>Batch Normalization</p> <ul style="list-style-type: none"> • Axis: -1 • Momentum: 0.99 • Epsilon: 0.001 <p>Dense Layer</p> <ul style="list-style-type: none"> • Units: 256 <p>Regularization</p> <ul style="list-style-type: none"> • Kernel Regularization: 12(0.016) • Activity Regularization: 11(0.006) • Bias Regularization: 11(0.006) <p>Activation</p> <ul style="list-style-type: none"> • Activation Function: ReLU
---	--

Dropout

- **Rate:** 0.4

Output Layer

- **Units:** 14 (softmax activation for multi-class classification)

Optimizer

- **Optimizer:** Adamax
- **Learning Rate:** 0.001

Callbacks

- **ReduceLROnPlateau:**
 - **Monitor:** "val_loss"
 - **Factor:** 0.4
 - **Patience:** 2
 - **Min LR:** 0.0
- **EarlyStopping:**
 - **Monitor:** "val_loss"
 - **Patience:** 2
 - **Restore Best Weights:** True

Training Parameters

- **Number of Epochs:** 5

	<pre> base_model=tf.keras.applications.EfficientNetV2B1(include_top=False, weights="imagenet", input_shape=(255,255,3), pooling='max') print('Created EfficientNetV2 B1 model') Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/efficientnet_v2/efficientnetv2-b1_notop.h5 8456008/28456008 [=====] - 2s 0us/step Created EfficientNetV2 B1 model base_model.trainable=True x=base_model.output x=BatchNormalization(axis=-1, momentum=0.99, epsilon=0.001)(x) x=Dense(256, kernel_regularizer=regularizers.l2(1e-05), activity_regularizer=regularizers.l1(1e-05), bias_regularizer=regularizers.l1(0.006), activation='relu')(x) x=Dropout(rate=.4, seed=123)(x) output=Dense(14, activation='softmax')(x) model=Model(inputs=base_model.input, outputs=output) model.compile(optimizer=Adamax(learning_rate=0.001), loss='categorical_crossentropy', metrics=['accuracy']) rlrmp=keras.callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.4, patience=2, verbose=1, mode='auto', min_delta=0.00001, cooldown=0, min_lr=0.0) estop=keras.callbacks.EarlyStopping(monitor='val_loss', min_delta=0, patience=2, verbose=1, mode='auto', baseline=None, restore_best_weights=True) callbacks=[rlrmp, estop] history=model.fit(x=train_gen, epochs=5, verbose=1, callbacks=callbacks, validation_data=val_gen, validation_steps=None, shuffle=True) </pre>
Vgg19	<p>Hyperparameter Tuning:</p> <p>Base Model</p> <ul style="list-style-type: none"> • Base Model: VGG19 (pretrained on ImageNet, used as a feature extractor) • Input Shape: (224, 224, 3) <p>Pooling</p> <ul style="list-style-type: none"> • Pooling: Global Average Pooling <p>Dense Layers</p> <ul style="list-style-type: none"> • Units: 1024 • Activation Function: ReLU <p>Output Layer</p> <ul style="list-style-type: none"> • Units: len(train_gen.class_indices)

- **Activation Function:** Softmax

Optimizer

- **Optimizer:** Adam
- **Learning Rate:** 0.001

Loss Function

- **Loss Function:** Categorical Crossentropy

Metrics

- **Metrics:** Accuracy

Freezing Layers

- **Freezing Layers:** Initially freeze all base model layers

```
from tensorflow.keras.applications import VGG19
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D
from tensorflow.keras.optimizers import Adam
base_model = VGG19(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(1024, activation='relu')(x)
predictions = Dense(len(train_gen.class_indices), activation='softmax')(x)

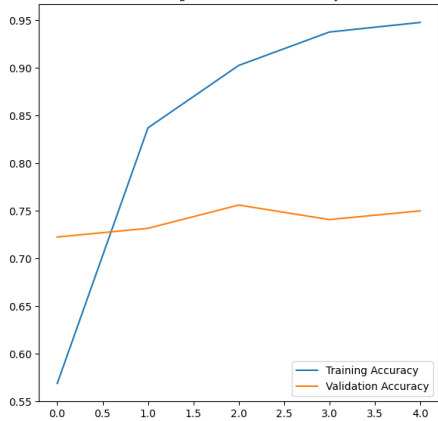
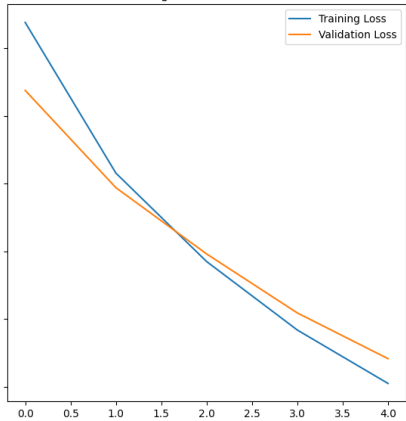
model = Model(inputs=base_model.input, outputs=predictions)

for layer in base_model.layers:
    layer.trainable = False

model.compile(optimizer=Adam(learning_rate=0.001),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.summary()
```


Final Model Selection Justification (2 Marks):

Final Model	Reasoning
	<p>EfficientNetV2B1:</p> <div style="display: flex; justify-content: space-around;">   </div> <ol style="list-style-type: none"> 1. Training Accuracy: <ul style="list-style-type: none"> ○ The training accuracy starts around 0.55 and increases rapidly, reaching around 0.95 by epoch 4. 2. Validation Accuracy: <ul style="list-style-type: none"> ○ The validation accuracy starts around 0.70 and remains relatively stable, fluctuating slightly but generally around 0.75 by epoch 4. 3. Training Loss: <ul style="list-style-type: none"> ○ The training loss decreases significantly from around 8 to below 3.
EfficientNetV2B1 Model Initialisation	<ul style="list-style-type: none"> ○ The training loss decreases significantly from around 8 to below 3.

4. Validation Loss:

- The validation loss also decreases from around 7 to slightly above 3.

VGG19:



1. Training Accuracy:

- The training accuracy starts around 0.4 and increases steadily, reaching around 0.9 by epoch 17.5.

2. Validation Accuracy:

- The validation accuracy starts around 0.4 and shows fluctuations, peaking around 0.7 but generally staying lower than the training accuracy.

3. Training Loss:

- The training loss decreases significantly from around 2 to nearly 0.2.

4. Validation Loss:

- The validation loss decreases initially from around 1.8 to around 1.2 but then fluctuates and increases slightly.

Comparison:

- **Training Accuracy:**

- EfficientNetV2B1 reaches higher training accuracy more quickly than VGG19.

- **Validation Accuracy:**

- EfficientNetV2B1 maintains a relatively stable validation accuracy around 0.75, while VGG19's validation accuracy fluctuates and generally stays lower.

- **Training Loss:**

- Both models show a significant decrease in training loss, but EfficientNetV2B1 starts with a higher value and decreases more sharply initially.

- **Validation Loss:**

- EfficientNetV2B1 has a more consistent decrease in validation loss, whereas VGG19's validation loss fluctuates after the initial decrease.

Conclusion:

EfficientNetV2B1 seems to perform better overall, with higher and more stable validation accuracy and a consistent decrease in validation loss compared to VGG19. While both models improve over time,

EfficientNetV2B1 shows quicker and more consistent improvements in the metrics provided.

```
#### EfficientNetV2B1 Model Initialisation

base_model=tf.keras.applications.EfficientNetV2B1(include_top=False, weights='imagenet',input_shape=(255,255,3), pooling='max')
print('Created EfficientNetV2 B1 model')

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/efficientnet_v2/efficientnetv2-b1_notop.h5
8456008/28456008 [=====] - 2s 0us/step
Created EfficientNetV2 B1 model

base_model.trainable=True
x=base_model.output
x=BatchNormalization(axis=-1, momentum=0.99, epsilon=0.001 )(x)
x = Dense(256, kernel_regularizer = regularizers.l2(l = 0.016),activity_regularizer=regularizers.l1(0.006),
      bias_regularizer=regularizers.l1(0.006) ,activation='relu')(x)
x=Dropout(rate=.4, seed=123)(x)
output=Dense(14, activation='softmax')(x)
model=Model(inputs=base_model.input, outputs=output)
model.compile(Adamax(learning_rate=0.001), loss='categorical_crossentropy', metrics=['accuracy'])
```

```
rlrmp=keras.callbacks.ReduceLROnPlateau(monitor="val_loss", factor=0.4, patience=2, verbose=1, mode="auto", min_delta=0.00001, cooldown=0, min_lr=0.0001)
earlystop=keras.callbacks.EarlyStopping(monitor="val_loss", min_delta=0, patience=2, verbose=1, mode="auto", baseline=None, restore_best_weights=True)
callbacks=[rlrmp, earlystop]

Model fitting

history=model.fit(x=train_gen, epochs=5, verbose=1, callbacks=callbacks, validation_data=val_gen,
                  validation_steps=None, shuffle=True)

Epoch 1/5
100/100 [=====] - 94s 423ms/step - loss: 8.3989 - accuracy: 0.5595 - val_loss: 7.3195 - val_accuracy: 0.7043 - lr: 0.0010
Epoch 2/5
100/100 [=====] - 33s 330ms/step - loss: 6.1665 - accuracy: 0.8233 - val_loss: 5.9759 - val_accuracy: 0.7134 - lr: 0.0010
Epoch 3/5
100/100 [=====] - 33s 334ms/step - loss: 4.8552 - accuracy: 0.9032 - val_loss: 4.9639 - val_accuracy: 0.7317 - lr: 0.0010
Epoch 4/5
100/100 [=====] - 34s 335ms/step - loss: 3.8550 - accuracy: 0.9258 - val_loss: 4.1046 - val_accuracy: 0.7256 - lr: 0.0010
Epoch 5/5
100/100 [=====] - 34s 337ms/step - loss: 3.0294 - accuracy: 0.9486 - val_loss: 3.4238 - val_accuracy: 0.7195 - lr: 0.0010
```